BusyLean: Naming Ever Larger Numbers in Lean

Research notes on constructive large numbers with formal comparisons

BusyLean contributors

October 2026

Abstract

We port Cody Roux's *name the biggest number* contest from Coq to Lean 4. In the contest every number must be a closed, computable definition, and every claim "my number is bigger" must be a kernel-checked theorem. We give an exact port of the incumbent Contender 5 (the largest numeral of a raw System T term of depth 42). We review and rebuild five submitted proposals, and we construct Contender 6, which combines their best ideas: an object language with W-types over every finite type, relativized to oracle families, and reflected along every tree that the language itself can define. Along the way we prove formally that one proposal's re-implementation of the old language agrees with the exact port. We then construct Contender 7: System F with Contender 6 as an oracle. Its strong normalization is proved inside Lean by Girard's reducibility candidates, which is possible because Lean's `Prop` is impredicative, and normalization runs as a certified total function. The next rungs of the ladder are System F_ω and the Calculus of Constructions.

Contents

1 The game **2**

2 How to compare unimaginably large numbers **2**

3 Contender 5, exactly **2**

4 Five proposals **3**

5 Contender 6 **3**
- 5.1 The object language . . . 3
- 5.2 Oracles for free . . . 3
- 5.3 Reflection along the language's own trees . . . 4
- 5.4 What is proved . . . 4

6 Engineering lessons **4**

7 Contender 7: impredicative polymorphism **5**
- 7.1 Why Lean can host System F . . . 5
- 7.2 The formalization . . . 5
- 7.3 From programs to numbers . . . 5
- 7.4 The contender . . . 5

8 Roadmap **6**

9 Status **6**

1 The game

Scott Aaronson's *Who can name the bigger number?* asks for the largest number one can describe concisely. Cody Roux's repository name-the-biggest-number adds two caveats that turn the game into one about *proof*. First, numbers must be *constructive*: computable in principle, with no Busy Beavers. Second, each new contender must come with a *formal proof* that it beats the previous one. The original contest uses Coq. Its last entry, Contender 5 by Samuel Gruetter, is the largest value of a raw STLC+NatRec term (Gödel's System T) of syntax depth at most 42.

The Lean edition. This repository keeps the contest but switches to Lean 4 (version 4.34.0). The rules (`README.md`) are: closed computable `Nat` definitions; a theorem `contender_(N-1)_lt_contender_N`; the axioms `propext` and `Quot.sound` allowed and `Classical.choice` forbidden; at most 15 s to check the definition and 60 s to check the proof, from a clean build; and no laziness or shenanigans. The axiom policy is deliberate. `Classical.choice` is the only one of Lean's three axioms that conflicts with computability, so it is the one that must go.

2 How to compare unimaginably large numbers

All numbers here are far too large to evaluate. The formal theorems compare them directly, but informal *size* is best discussed with the *fast-growing hierarchy* f_α indexed by ordinals: $f_0(n) = n+1$, $f_{\alpha+1}(n) = f_\alpha^n(n)$, and $f_\lambda(n) = f_{\lambda[n]}(n)$ at limits. A total language L whose programs are all provably terminating in a theory T has its bounded-size maxima growing roughly like $f_{|T|}$, where $|T|$ is the proof-theoretic ordinal of T. Two heuristics organize everything below.

1. **The base language dominates.** System T is calibrated by ε_0. T plus one type of countably branching trees, with recursion into all finite types, is calibrated by the Bachmann–Howard ordinal $\mathrm{BH} = \psi(\varepsilon_{\Omega+1})$. With tree classes at every finite level the calibration is $\psi(\Omega_\omega)$, the ordinal of Π^1_1-CA_0.

2. **Reflection multiplies.** Relativizing a language L to an oracle that grows like f_β yields growth like $f_{\beta+|L|}$. Iterating this along an index of height γ yields $f_{|L|\cdot\gamma}$. Explicit ordinal notations without collapsing functions stay below BH, so a stack of reflections over a weak base cannot catch up with a stronger base.

These are informal calibrations, not theorems; the formal content of every contender is a strict inequality checked by Lean's kernel.

3 Contender 5, exactly

`BusyLean/Contender5.lean` ports `Contender.v` one definition at a time. The port keeps every quirk that can affect the value: raw ill-typed terms are interpreted with `error` defaults; variables use reversed de Bruijn indices; lambdas carry both annotations and cast their body; arrow casts recurse only when `type_eqb` succeeds; the enumeration order is the original one, including right-nested `++`;

and the winner is chosen by `maxBy`'s strict update, starting from `t0`. A smoke test reproduces the reference table: $0, 2, 10, 146, 24976$ terms and maxima $0, 0, 1, 2, 3$ at depths 0–4.

Agreement of an independent port. Proposal P5 re-implemented the old language *inside* its larger W-language, with W-aware casts and a plain maximum instead of `maxBy`, and asserted without proof that the result is Contender 5. We prove it:

Theorem 1 (`PortAgreement.largest_eq_oldMaximum`)**.** *For every depth n,*

$$\mathit{C5.largestSTLCNatRecNatOfDepth}\ n\ =\ \mathit{BigNumber.oldMaximum}\ n.$$

The proof is a logical relation between the two interpreters: equality at $\mathbb{N}$, preservation at arrows. That avoids function extensionality. Casts in both directions preserve the relation, including the error branches. The fundamental lemma covers reversed lookup, binder casts and natural recursion. P5's enumeration is the translated image of the original, in the same order, and the `maxBy` winner attains the maximum.

4 Five proposals

Five proposals arrived as archives; `proposals/REVIEW.md` has the details. All were rebuilt on Lean 4.34.0, and the two that had never been compiled (P4, P5) needed small repairs, recorded in separate commits.

	base language	amplifier	informal size
P1	System T	FGH over explicit Veblen-style trees	$f_{\varepsilon_0+\alpha}$, $\alpha < \mathrm{BH}$
P2	T + Brouwer trees + oracles	forest reflection on explicit trees	$f_{\mathrm{BH}\cdot\alpha}$
P3	T + unary oracle	reflection on explicit trees	$f_{\varepsilon_0\cdot\alpha}$
P4	T + Brouwer trees + oracle	ramified reflection	$f_{\mathrm{BH}\cdot\alpha}$
P5	T + W-types at all finite types	two diagonals	$f_{\psi(\Omega_\omega)+\omega^2}$

P5 has by far the strongest base and the weakest amplifier, which suggests combining it with the others' amplifiers.

5 Contender 6

5.1 The object language

The language is P5's. Types are generated by $A, B ::= \mathbb{N} \mid A \to B \mid W\,A$, where $W\,A$ is the type of well-founded trees branching over A. Terms are raw (ill-typed terms default), with natural recursion and a W-recursor into every type. Because A ranges over all types, the language has tree classes of every finite level, $W\mathbb{N}$, $W(W\mathbb{N})$, …, and its calibration is $|L| \approx \psi(\Omega_\omega)$.

5.2 Oracles for free

P5's interpreter resolves variables by oldest-first levels. Evaluating a program in the one-element environment $[\langle \mathbb{N} \to \mathbb{N} \to \mathbb{N}, fs\rangle]$ therefore makes `var 0` denote the oracle family fs everywhere, even under binders. No new syntax is needed, and the smoke test confirms it: `var0 3 4` evaluates to $fs(3,4)$, and so does `(fun y => var0 y y) 5` with the family applied at $(5,5)$.

5.3 Reflection along the language's own trees

```
def reflect (fs : Family) (n : ℕ) : ℕ :=
  n + 1 + max (prefixMax (fun i => fs i n) n)
              (maximum (evalWith fs) (allTerms n))
def stage (seed : Family) : WTree ℕ → ℕ → ℕ
  | .leaf   => reflect seed
  | .node f => reflect (fun i => stage seed (f i))
def treesOfDepth (k : ℕ) : WTree ℕ :=
  .node fun i => defTree (nthD (allTerms k) i .zero)
def index : WTree ℕ := .node treesOfDepth
def contender_6 : ℕ := stage p5Family index 42
```

At a node the oracle family *is* the family of child stages, so a program may call any child at a computed index (P2's forest reflection). The new ingredient is `index`: its k-th child joins every $W\mathbb{N}$ tree denoted by a program of depth $\leq k$. The hierarchy therefore runs along the language's own ordinals, all of $|L|$, instead of an explicit notation system below BH. By the multiplication heuristic, Contender 6 is about $f_{|L|\cdot|L|}(42)$.

5.4 What is proved

- `contender_5_lt_contender_6`: the exact port of Contender 5 is below Contender 6.
- `p5_lt_contender_6` and `p5Family_lt_finalFamily`: P5 is beaten at every input.
- `defined_tree_eventually_lt`: every tree the language defines is eventually dominated.
- `program_lt_finalFamily`: every program of depth $\leq n$, run with the final oracles, is exceeded.

All of them depend only on `propext`. The definitions check in about 5 s and the proofs in about 6 s. A machine-checked audit shows that the definition never mentions `contender_5` or its winner selection. It also reports, candidly, that P5's seed contains the old language's maximum as a function of depth.

6 Engineering lessons

- **Toolchain drift changes axiom lists.** In Lean 4.34 a wildcard match arm compiles to a `_sparseCasesOn` helper built on `Nat.land`. `Nat.bitwise`'s termination proof uses `propext`, so "axiom-free on 4.19" did not survive the upgrade. The same happens to core `xs[i]?`. Exhaustive matches avoid it.
- **Tactics carry axioms.** `omega`'s certificates use `funext`, and so `Quot.sound`. Two calls were replaced by explicit lemma chains.
- **Never let the unifier normalize a contender.** Writing `iterate f 41 (f 42)` against `iterate f 42 42` made Lean try to evaluate the number. Generalizing the statement first, so that only $42 = 41 + 1$ is checked, fixes it.
- **Prove inequalities by inclusion, not monotonicity.** Reflection is not monotone in its oracle, because a program may test oracle values. Every comparison therefore exhibits an explicit occurrence: a visible oracle index or an enumerated program.

7 Contender 7: impredicative polymorphism

Contender 6 lives below the next collapsing level: everything it does is predicative. The next qualitative jump is *impredicative polymorphism*. The numeric functions definable in Girard's System F on Church numerals are exactly those provably total in second-order arithmetic Z_2, and $\psi(\Omega_\omega)$ is merely the ordinal of Π^1_1-CA$_0$, a small fragment of Z_2.

7.1 Why Lean can host System F

System F has no set-theoretic model in Lean's predicative `Type`: a model would have to interpret $\forall X.A$ as a product over *all* types, including itself. Girard's normalization proof needs much less. It interprets types as *reducibility candidates*, predicates on terms, and quantifies over all candidates. In Lean a candidate is a structure whose main field is $P : \mathrm{Tm} \to \mathrm{Prop}$. A statement "$\forall C : \mathrm{Cand}, \ldots$" is again a proposition, because `Prop` is impredicative. The normalization proof therefore goes through, even though there is no semantic model.

7.2 The formalization

1. **Untyped core** (`SystemF/Syntax`, `Reduction`): de Bruijn terms with an oracle constant `orc`; parallel substitution with its composition algebra; full β-reduction plus the δ-rule $\texttt{orc}\,\overline{n} \to \overline{f(n)}$ on canonical Church numerals; SN as accessibility.

2. **Types** (`Types`): de Bruijn System F types and Curry-style typing of *erased* terms. A flag enables the oracle rule; numerals are typed without it.

3. **Candidates** (`Candidates`): CR1–CR3, where *neutral* means neither a λ nor the oracle (`orc a` can reduce at the head); the arrow construction; the intersection $\forall$; the interpretation $[\![A]\!]_\rho$; its renaming and substitution lemmas, proved as *equalities* of candidates via funext, propext and proof irrelevance; β-closure for λ; and adequacy. The oracle is justified afterwards. Numerals lie in every interpretation of `nat` (adequacy without the oracle rule), so the δ-reduct of $\texttt{orc}\,a$ stays inside the candidate.

Theorem 2 (`F.sn_of_typed`)**.** *Every term typable in System F with the oracle is strongly normalizing under $\beta\delta$, for every oracle function f. The proof uses only* `propext` *and* `Quot.sound`*.*

7.3 From programs to numbers

Programs are Church-style terms (annotated λ, explicit type abstraction and application), so type inference is decidable. `infer_sound` states that inference implies the Curry typing of the erased term. A program's value is computed as follows: check that it is closed of type $\texttt{nat} = \forall X.(X \to X) \to X \to X$; erase it; normalize it by the normal-order step function `nstep`, iterated by well-founded recursion on the SN certificate (the compiler runs this as an ordinary loop); and read the Church numeral. A smoke test shows the pipeline computing: $2^3 \mapsto 8$ and $3^{2^2} \mapsto 81$ by polymorphic instantiation, oracle calls $\texttt{orc}\,\overline{4} \mapsto 40$ for $f(n) = 10n$, and 0 for ill-typed programs.

7.4 The contender

```
def reflectF (f : ℕ → ℕ) (n : ℕ) : ℕ :=
  n + 1 + max (f n) (maximum (fvalue f) (ctmsUpTo n))
def F7 (n : ℕ) : ℕ := reflectF finalFamily n
```

```
def contender_7 : ℕ := F7 (F7 42)
```

The oracle is Contender 6's family. System F programs can call it and amplify it with second-order strength. The outer `F7` runs System F at a syntax budget already larger than Contender 6. Proved: `contender_6_lt_contender_7`, `finalFamily_lt_F7` (at every input), and `fvalue_lt_contender_7` (every System F program of depth $\leq$ `F7 42` is exceeded). The audit confirms that the constant `contender_6` itself never occurs in the definition; only Contender 6's family occurs, as the oracle.

Size. By the heuristics of Section 2, Contender 7 is about $f_{|\mathsf{Z}_2|}$ applied at an input beyond Contender 6. Here $|\mathsf{Z}_2|$ is the proof-theoretic ordinal of second-order arithmetic. Nobody has computed it, but it lies far beyond every ordinal named in this article.

8 Roadmap

System $\mathbf{F}_\omega$. Adding type operators (kinds $\kappa ::= * \mid \kappa \to \kappa$) gives the strength of higher-order arithmetic. The candidate method extends: a type constructor of kind $\kappa \to \kappa'$ is interpreted as a function between candidate spaces, defined by recursion on kinds. Type equality now includes β-conversion of type operators, so Church-style checking needs type normalization, which is itself the normalization of the simply typed λ-calculus.

The Calculus of Constructions. Dependent types on top of F_ω is the system behind Loader's number, the largest well-known computable googologism. Its normalization proof (Coquand, Geuvers) also uses candidates, now indexed by the dependent structure.

Amplifiers. The reflection machinery of Contender 6 applies to any total language with an oracle. A forest reflection along trees definable in System F would add a multiplicative factor. That is modest compared to the next rung of the ladder, which remains the priority.

9 Status

Contender	idea	informal size
5 (Coq)	System T, depth 42	$f_{<\varepsilon_0}(42)$
6 (Lean)	W-language, reflected along its own trees	$f_{\psi(\Omega_\omega)^2}(42)$
7 (Lean)	System F with Contender 6 as oracle	$f_{\lvert\mathsf{Z}_2\rvert}(>C_6)$
8 (planned)	System F_ω	higher-order arithmetic